

Chapter 03

Configuring the Shell



Introduction

- One key component of the Bash shell is shell *variables*.
- Variables store vital system information and modify the behavior of the Bash shell, as well as many commands.
- The `PATH` variable affects how commands are executed and how other variables affect your ability to use the history of your commands.
- Initialization files make shell variables *persistent*, so they will be created each time you log into the system.

Shell Variables

- A *variable* is a name or identifier that can be assigned a value.
- To assign a value to a variable, type the name of the variable immediately followed by the equal sign = character and then the value.

```
name="value"
```

- Variable names should start with a letter (alpha character) or underscore and contain only letters, numbers and the underscore character. For example:
 - `_a=1`
 - `_1=a`
 - `LONG_VARIABLE='OK'`
 - `Name='Jose Romero'`

2019 © Network Development Group Inc.



Local Environment Variables

- A *local variable* is only available to the shell in which it was created.
- An *environment variable* is available to the shell in which it was created, and all other commands/programs started by the shell.
- To set the value of a variable, use the following assignment expression.

```
variable=value
```

```
sysadmin@localhost:~$ name='julie'
sysadmin@localhost:~$ echo $name
julie
```

- To create an environment variable, use the `export` command.

```
sysadmin@localhost:~$ export JOB=engineer
```

2019 © Network Development Group Inc.



Unsetting Variables

- If you create a variable and then no longer want that variable to be defined, use the `unset` command to delete it:

```
unset VARIABLE
```

Warning

Do not unset critical system variables like the `PATH` variable, as this may lead to a malfunctioning environment.

Displaying Variables

- There are several ways to display the values of variables.
- The `set` command will display all variables (local and environment).
- To display only environment variables, you can use several commands that provide nearly the same output:
 - `env`
 - `declare -x`
 - `typeset -x`
 - `export -p`
- To display the value of a specific variable, use the `echo` command with the name of the variable prefixed by the `$` (dollar sign). For example:

```
sysadmin@localhost:~$ echo $PATH
```

PATH Variable

- The `PATH` variable contains a list of directories that are used to search for commands entered by the user.
- The `PATH` directories are searched for an executable file that matches the command name.
- The following example displays a typical `PATH` variable:

```
sysadmin@localhost:~$ echo $PATH
/home/sysadmin/bin:/usr/local/sbin:/usr/local/bin:/usr/sbin:/usr/bin:/sbin:/bin:/usr/games:/usr/local/games
```

PATH Variable

- To execute commands that are not contained in the directories that are listed in the `PATH` variable, several options exist:
 - Type the *absolute path* to the command.
 - Use the *relative path* to the command.
 - The `PATH` variable can be set to include the directory where the command is located.
 - Copy command to a directory that is listed in the `PATH` variable.
- An *absolute path* specifies the location of a file or directory from the top-level directory (i.e. `/usr/bin/ls`).
- A *relative path* specifies the location of a file or directory relative to the current directory (i.e. `test/newfile`).

Initialization Files

- Initialization files set the value of variables, create aliases and functions, and execute other commands that are useful in starting the shell.
- There are two types of initialization files:
 - Global initialization files - affect all users on the system.
 - Local initialization files - specific to an individual user.
- BASH initialization files include:
 - `/etc/profile`
 - `~/.bash_profile`, `~/.bash_login`, `~/.profile`
 - `~/.bashrc`
 - `/etc/bashrc`

2019 © Network Development Group Inc.



Modifying Initialization Files

- The way a user's shell operates can be changed by modifying that user's initialization files.
- In some distributions, the default `~/.bash_profile` file contains lines that customize the `PATH` environment variable:

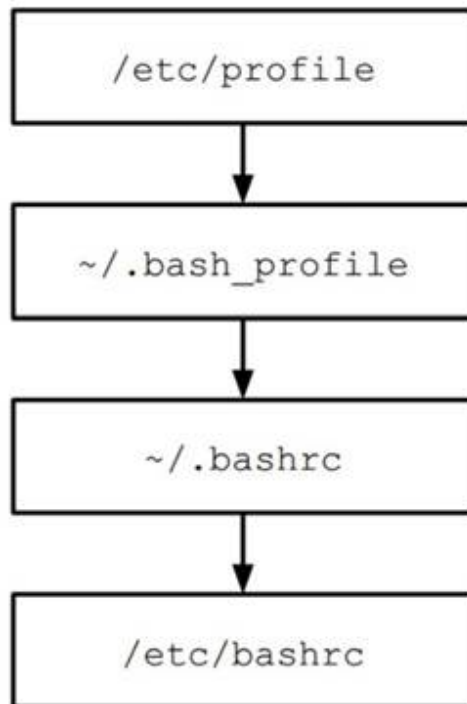
```
PATH=$PATH:$HOME/bin
export PATH
```

- The first line sets the `PATH` variable to the existing value with the addition of the `bin` subdirectory of the user's home directory.
- The second line converts the local `PATH` variable into an environment variable.

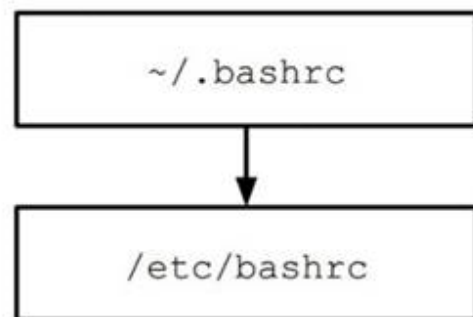
2019 © Network Development Group Inc.



Login Shell



Interactive Shell



BASH Exit Scripts

- The Bash shell may execute one or more files upon exiting.
- These files are used for "cleaning up" as the user exits the shell.
- The following exit files may exist:
 - `~/.bash_logout`
 - `/etc/bash_logout`

Command History

- The `~/.bash_history` file contains a history of the commands that a user has executed within the Bash shell.
- There are several ways that this command history is advantageous to the user:
 - The **Up ↑** and **Down ↓ Arrow Keys** can be used to review your history and select a previous command to execute again.
 - Select a previous command and modify it before executing it.
 - Press **Ctrl+R** and then begin typing a portion of a previous command to do a reverse search through history.
 - Execute a command again, based upon a number that is associated with the command.

Configuring the history Command

- When the shell is closed, commands in the history list are stored in `~/.bash_history`, also called the *history file*.
- The `HISTFILESIZE` variable will determine how many commands to write to this file.
- To store the history commands in a different file, edit the value of the `HISTFILE` variable.
- The `HISTCONTROL` variable can be set to different features such as ignoring spaces or duplicate commands.
- The `HISTIGNORE` variable can also be used to ignore commonly used commands.

2019 © Network Development Group Inc.



Using the history Command

- The `history` command can be used to re-execute previously executed commands.

```
sysadmin@localhost:~$ history
1 ls
2 cd test
3 cat alpha.txt
4 ls -l
5 cd ..
```

- The most common options for the `history` command are:
 - `-c` = Clear list
 - `-r` = Read the history file and replace the current history
 - `-w` = Write the current history list to the history file



Execute Previous Commands

- The **!** exclamation mark is a special character that indicates the execution of a command within the history list.
- The following are some examples of using the exclamation **!** character:
 - **!!** - Repeat the last command
 - **!-4** - Execute command that was run four commands ago
 - **!55** - Execute command number 55
 - **!to** - Execute the last command that starts with **to**
 - **!?bob** - Execute the last command that contained **bob**

MODULO 8

Grep permette di cercare all'interno di un testo delle parole specifiche (grep 'uva' file.txt, grep '^a' file.txt, grep 'a\$' file.txt)

```
cate@cate-virtual-machine:~$ touch frutta
cate@cate-virtual-machine:~$ nano frutta
cate@cate-virtual-machine:~$ cat frutta | grep 'uva'
uva
cate@cate-virtual-machine:~$ grep '^a' frutta
albicocca
avocado
cate@cate-virtual-machine:~$ ls
archivio  Documents  elenco1      example  file1  fileabc  fileMaster  frutta  Pictures  risultato  snap      verde
Desktop  Downloads  elenco_file.txt  file    file2  fileL    fileN       Music   Public    script     Templates Videos
cate@cate-virtual-machine:~$ nano frutta
cate@cate-virtual-machine:~$ cat frutta
mela
banana
albicocca
uva
avocado
cate@cate-virtual-machine:~$ grep 'uva' frutta
uva
cate@cate-virtual-machine:~$ grep 'a$' frutta
mela
banan
albicoc
uva
cate@cate-virtual-machine:~$ ^ = inizio parola, $ = fine parola
```

Chapter 14

Filesystem Links



Introduction

- A link is a way to make a file's data accessible via more than one filename.
- There are two types of links available:
 - Soft links
 - Hard links

Soft Links

- Also called *Symbolic Links*.
- Soft links are a file type designed to point to another file using a path name.
- Soft links are distinguishable by their file type.
- For example, a detailed listing of the `/bin/systemd` file shows that it is a symbolic link:

```
sysadmin@localhost:~$ ls -l /bin/systemd
lrwxrwxrwx 1 root root 20 Feb 28 21:03 /bin/systemd -> /lib/systemd/systemd
```

`lrwxrwxrwx 1 root root 20 Feb 28 21:03 /bin/systemd -> /lib/systemd/systemd`

Soft Links

- Soft links are excellent for creating "shortcuts".
- When a system file is moved to another location by the developers, soft links are created to make it easier for administrators and users to find the new location.

Soft Links

- To create a soft link file, use the `ln` command with the `-s` option.

```
ln -s <target> <link_name>
```

```
sysadmin@localhost:~$ ln -s file1.txt file2.txt
sysadmin@localhost:~$ ls -l file*
-rw-rw-r-- 1 sysadmin sysadmin 0 May 9 02:48 file1.txt
lrwxrwxrwx 1 sysadmin sysadmin 9 May 9 02:49 file2.txt -> file1.txt
```

- Soft link files *do not* increase the *link count* number associated with a regular file.

```
-rw-r--r-- 1 sysadmin sysadmin May 9 14:39 file1.txt
lrwxrwxrwx 1 sysadmin sysadmin May 9 14:39 file2.txt -> file1.txt
```

Hard Links

- Hard links are two or more files that share the same inode number.
 - Inode number: index node - contains file info including location on the disk.
- Hard links are exactly identical to the original in every way except the file name
- Hard links are created using the `ln` command without using the `-s` option

```
ln <target> <link_name>
```


Hard Links

- Hard links identified using `ls -l`.

```
sysadmin@localhost:~$ ls -l profile.txt
-rw-r--r-- 2 sysadmin sysadmin 110 Apr 24 16:24 profile.txt
```

- The link count (2) indicates how many hard links there are to the file.
- Since hard links share the same inode, they will have the same inode number.
- The `ls -li` option can be helpful to validate that the files are sharing an inode:

```
sysadmin@localhost:~$ ls -li profile.txt myprofile.txt
95813671 -rw-r--r-- 2 sysadmin sysadmin 110 Apr 24 16:24 myprofile.txt
95813671 -rw-r--r-- 2 sysadmin sysadmin 110 Apr 24 16:24 profile.txt
```

2019 © Network Development Group Inc.



Soft Links vs Hard Links

- Advantages of Hard links:
 - Hard linked files are indistinguishable by programs from regular files
 - If files are hard linked then they are always contained within one filesystem
 - It is easy to find files that are hard linked
 - Removing hard links doesn't remove the actual data unless you remove all of the hard links

2019 © Network Development Group Inc.



Soft Links vs Hard Links

- Advantages of Soft links:
 - Soft links can be made to a directory file; hard links can not.
 - Soft links can be made from a file on one filesystem to a file on another filesystem; hard links can not.
 - Soft links are very visual because the output of the `ls -l` command displays which file the soft link is pointing to.

PRO:

Ottimizzazione dello spazio e come e cosa volete creare

Chapter 11

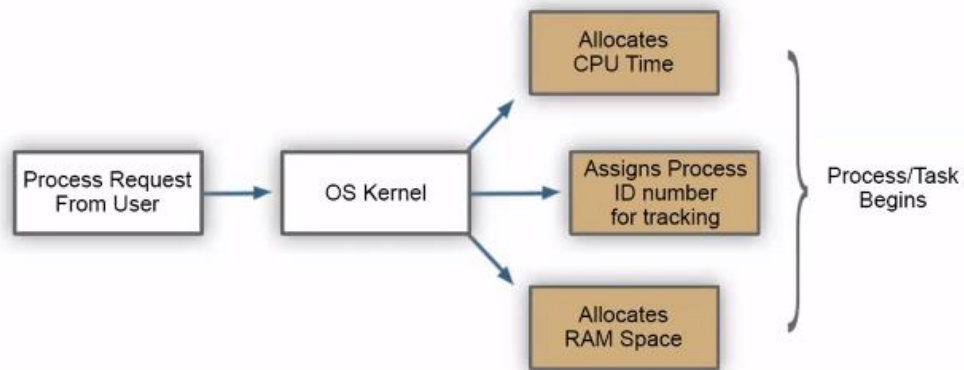
Managing Processes



Introduction

- Managing processes covers:
 - How to run processes in the background or foreground.
 - How to make a process switch between the background and foreground.
 - How to control processes by sending them signals using the `kill` command.
 - Techniques for monitoring the resources that a process is using.
 - How to control the priority of processes.
 - Utilities that allow multiple processes inside a single shell.

Process Initiation and Tracking



Process Control

- Running a command results in something called a *process*.
- Linux manages tasks using processes.
- A process can start a subprocess and form a parent/child relationship.
- Users can only control their processes.
- Root can control all system and user processes

Listing Processes

- The `ps` (process status) command lists running processes
- By itself, lists processes running in the current terminal:

```
ps [OPTION]...
```

```
sysadmin@localhost:~$ ps
PID TTY          TIME CMD
 80 ?            00:00:00 bash
 94 ?            00:00:00 ps
```

PID	Process identifier unique to each process.
TTY	Name of the terminal or pseudo-terminal where the process is running.
TIME	Total processor time used by the process.
CMD	Command that started the process.

Listing Processes

- Use the `-e` (every) and `-f` (full) options with the `ps` command to list all processes on the system:

```
ps -ef
```

```
sysadmin@localhost:~$ ps -ef
UID          PID  PPID  C  STIME TTY          TIME CMD
root           1      0  0   17:16 ?            00:00:00 /sbin/?? /init
syslog        33      1  0   17:16 ?            00:00:00 /usr/sbin/rsyslogd
root          38      1  0   17:16 ?            00:00:00 /usr/sbin/cron
root          40      1  0   17:16 ?            00:00:00 /usr/sbin/sshd
bind          57      1  0   17:16 ?            00:00:00 /usr/sbin/named -u bind
root          70      1  0   17:16 ?            00:00:00 /bin/login -f
sysadmin      80     70  0   17:16 ?            00:00:00 -bash
sysadmin      96     80  0   17:26 ?            00:00:00
```


Searching For Processes

- The `pgrep` command looks for the specified names or other criteria of processes.
- The `-i` option makes the `pgrep` command case insensitive.
- For example, to find all instances of the `sshd` command running on the system:

```
sysadmin@localhost:~$ pgrep -i sshd
15
```

- The `-l` option can be used with the `pgrep` command to list the process name along with the PID
- The `-u` option can be used to search for processes owned by a specific user.

Watching Processes

- The `watch` command can be used to monitor recurring processes.
- The `watch` command can be used with the `ps` command to monitor running processes in the shell:

```
sysadmin@localhost:~$ watch ps aux
```

```
Every 2.0s: ps aux localhost: Fri Mar 29 17:47:56
2019
```

USER	PID	%CPU	%MEM	VSZ	RSS	TTY	STAT	START	TIME	COMMAND
root	1	0.0	0.0	18376	3048	?	Ss	17:44	0:00	/bin/bash /init
syslog	9	0.0	0.0	191328	3760	?	Ssl	17:44	0:00	/usr/sbin/rsyslogd

- The `watch` command's default interval is two seconds.
- To change the interval, use the `-n` option.

Executing Multiple Commands

- It can be useful to execute two or more commands in a single command line.
- By using the semicolon ; character as a delimiter between commands, a user can type multiple commands on one command line:

```
COMMAND;COMMAND[;COMMAND]...
```

- For example, to create an alias called `welcome` which outputs the current user, the date, and the current directory listing, execute the following command:

```
sysadmin@localhost:~$ alias welcome="whoami;date;ls"
sysadmin@localhost:~$ welcome
Sysadmin
Tue Mar 26 21:07:10 UTC 2019
Desktop Documents Downloads Music Pictures Public Templates Videos
```

Foreground Processes

- A foreground process is one that prevents the user from using the shell until the process is complete.
- When one process starts another, the first process referred to as the *parent* process and the new process is called a *child* process.

Background Processes

- When executed in the background, a child process releases control back to the parent process
- To have a command execute as a background process, add the ampersand & character after the command:

```
sleep 3 &
```

```
sysadmin@localhost:~$ sleep 3 &  
[1] 85
```

- When executing commands in the background, it outputs a *job number* followed by a space and then the *process identification number (PID)*.

2019 © Network Development Group Inc.



Managing Jobs

- The `jobs` command displays background jobs:

```
sysadmin@localhost:~$ jobs  
[1]-  Stopped      sleep 1000  
[2]+  Stopped      sleep 2000
```

- The `fg` (foreground) and `bg` (background) commands provide the ability to multi-task:

```
sysadmin@localhost:~$ fg 2  
sleep 2000
```

```
sysadmin@localhost:~$ bg 1  
[1]-  sleep 1000 &
```

2019 © Network Development Group Inc.



Moving Processes

- A command running in the foreground can be paused using Ctrl+Z:

```
sysadmin@localhost:~$ sleep 1000
^Z
[1]+  Stopped sleep 1000
```

- To put the paused command in the background, execute the **bg** command.

```
sysadmin@localhost:~$ bg
[1]+ sleep 1000 &
```

- A command that has been paused or sent to the background can be returned to the foreground using the **fg** command.

```
sysadmin@localhost:~$ fg
Sleep 1000
```

2019 © Network Development Group Inc.



Signals

- A signal is a message that is sent to a process to tell it to take some sort of action, such as stop, restart, or pause.
- Some signals can be sent to processes by simple keyboard combinations:

CTRL+z

CTRL+c

2019 © Network Development Group Inc.



Signals

- The `bg` command sends a process a signal to execute in the background.
- To list all signals, use the `kill -l` command.
- To send a process a signal, use the `kill` command followed by the PID# or Job# (prefixed with the percent % sign):

```
sysadmin@localhost:~$ kill 2901
```

```
sysadmin@localhost:~$ kill %1
```

```
sysadmin@localhost:~$ kill -p 2901
```

2019 © Network Development Group Inc.



Force Kill

- If other signals have failed to end a process, use the `SIGKILL` signal to force the process to end:

```
sysadmin@localhost:~$ kill -9 2901
```

```
sysadmin@localhost:~$ kill -KILL %1
```

```
sysadmin@localhost:~$ kill -SIGKILL -p 2901
```

2019 © Network Development Group Inc.



Other Signal Commands

- There are other commands that send processes signals such, as the `killall` and `pkill` commands
- They are useful to stop many processes at once
- To stop all processes owned by a user:

```
sysadmin@localhost:~$ killall -u bob
```

The HUP Signal

- When a user logs off the system, all processes that are owned by that user are automatically sent the Hang-up signal (SIGHUP)
- Typically, this signal causes those processes to end
- To have a process ignore HUP signals:

```
sysadmin@localhost:~$ nohup myjob.sh &
```

Process Priority

- Not all processes have the same access to the CPU
- A user of can influence the priority that will be assigned to a process by setting a *nice*ness value
- The higher you set the niceness value, the lower the priority that will be assigned to a process
- Highest: -20 Default: 0 Lowest: 19

Process Priority

- To set an initial niceness of a command, use the *nice* command:

```
sysadmin@localhost:~$ nice -n 19 cat /dev/zero > /dev/null
```

- To adjust the niceness of an existing process, use the *renice* command
- Only the root user can adjust a nice value to below 0 or lower than current value

Monitoring Processes

- The `top` command provides the ability to monitor processes in real-time, as well as manage the processes.

```
sysadmin@localhost:~$ top
```

- Monitors processes in real time using the interactive keys in `top`:
 - Press `h` to see all available options
 - Press `r` for renice
 - Press `q` to quit

Monitoring Processes

```
top - 16:47:34 up 51 days, 2:12, 1 user, load average: 1.37, 1.56, 1.49
Tasks: 13 total, 4 running, 9 sleeping, 0 stopped, 0 zombie
%Cpu(s): 0.3 us, 37.2 sy, 0.2 ni, 62.1 id, 0.0 wa, 0.1 hi, 0.0 si, 0.0 st
KiB Mem: 16438128 total, 13108516 used, 3329612 free, 4276 buffers
KiB Swap: 0 total, 0 used, 0 free. 9808716 cached Mem
```

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
164	root	20	0	4364	696	616	R	87.4	0.1	1:23.32	cat
165	root	30	10	4364	696	620	T	9.3	0.1	0:49.13	cat
166	root	39	19	4364	772	696	T	1.7	0.1	0:41.75	cat
1	root	20	0	17960	2972	2724	S	0.0	0.0	0:00.02	init
33	syslog	20	0	255844	2728	2296	S	0.0	0.0	0:00.03	rsyslogd
38	root	20	0	23656	2288	2076	S	0.0	0.0	0:00.00	cron
40	root	20	0	61364	3124	2444	S	0.0	0.0	0:00.00	sshd
57	bind	20	0	689640	29580	5328	S	0.0	0.2	0:00.13	named
70	root	20	0	63132	2900	2452	S	0.0	0.0	0:00.00	login
80	sysadmin	20	0	18176	3384	2896	S	0.0	0.0	0:00.04	bash
151	root	20	0	46628	2708	2360	S	0.0	0.0	0:00.01	su
152	root	20	0	18180	3388	2896	S	0.0	0.0	0:00.01	bash
167	root	20	0	19860	2452	2124	R	0.0	0.0	0:00.00	top

Monitoring the System

- The `uptime` command displays:
 - The current time
 - The amount of time the system has been running
 - The number of users who are currently logged in
 - The load averages during the past one, five and fifteen minute.

```
sysadmin@localhost:~$ uptime
18:24:58 up 5 days, 10:43, 1 user, load average: 0.08, 0.03, 0.05
```